

Comp 6321 - Machine Learning - Project report addendum

Federico O'Reilly Regueiro

December 20th, 2016

1 Some corrections and observations

At the time of submission there was still a bug in the code that greatly hampered performance and thus there had been little chance of simulating enough matches to see a real trend in the network. The bug has now been fixed by changing the way the gradient application functions are appended to the graph (init function starting at line 123), how these gradient applications are called (line 68) and how the gradients are reset (line 48 in the code in section 2.a. The original problem was that every epoch, an operation was being added to the computation graph, quickly consuming system memory. Now the operation is added only once to the graph and invoked by iterating a list of operation handles.

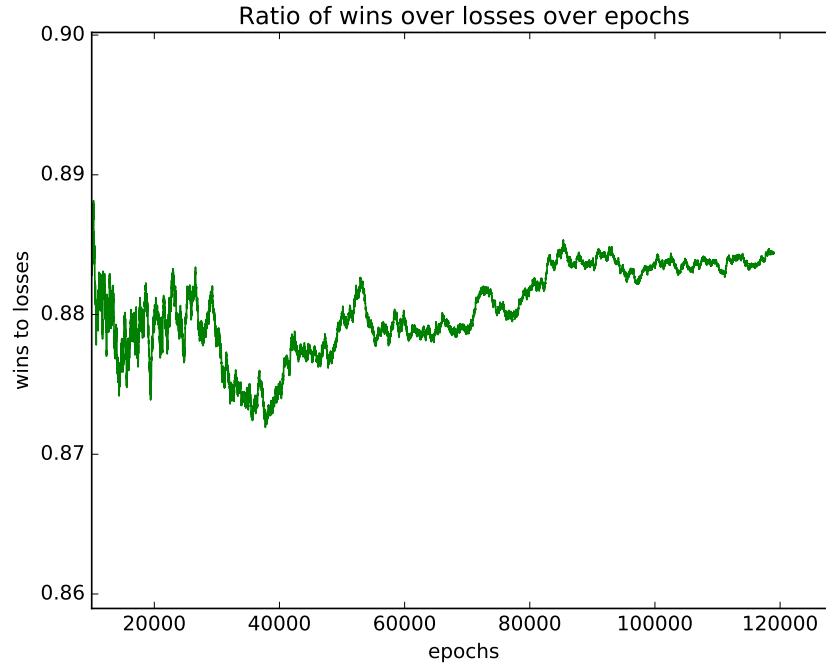


Figure 1: Ratio of wins and losses per epoch, we can see that the ratio is slowly improving but still under the expected ratio 1.0 for two random players.

This has allowed me to run the code for a sufficient number of epochs [4] to start seeing a trend in the network's performance as can be seen in figure ?? . Performance is still far from what I expected and improvement in performance is rather slow. This could be due to a number of reasons. I explore the possible reasons and strategies to be taken to minimize them:

- The network is too big: unlikely since some of the papers [2] [1] [3] report reasonable performance with larger networks.
- The learning rate is too low: possibly, Tesauro [4] reports reasonable performance with a learning rate of 0.1 after 125,000 iterations. The learning rate in the present implementation is one third of that.
- The nature of Othello, where the board configuration and particularly the material count can so quickly change, makes it difficult for the model to converge and it will simply take longer than Backgammon. Leouski and Utgoff [3] indicate similar performance with a NN based on board symmetry and only report better performance with a three-network model where responsibilities for opening, middle and end-game are spread across the three models.

- Either λ or γ are improperly tuned: different λ or γ settings might greatly impact the learning of the model. I have used Tesauro's reported values but given the aforementioned particularities of Othello, training might require different γ or λ
- The model from which we started is inadequate: gradient descent converges towards local minima and the model from which we started might have been close to a shallow local minimum, meaning that the model will never improve much. Try again with a different initial model whose random weights might be closer to a steeper minimum.
- The preprocessing performed, introducing domain-knowledge, is actually of little use to the net. Maybe the choice of rows, columns, 3×3 windows and diagonals is a good human heuristic but harder to learn from for a NN. Try a raw board input or another sort of pre-processing.

1.a Repository

I will most likely continue to work on this code, which has found its home at <http://github.com/friketrike/ML-AI-proj-2016> .

2 Appendix i - Code fixes

The memory leak was fixed by avoiding the addition of operations after initialization as described in section 1.

2.a ANN

```
1 import tensorflow as tf
2 import numpy as np
3
4 class othello_net():
5     @staticmethod
6     def weight_variables(shape, name=None):
7         initial = tf.random_normal(shape, stddev=0.1, name=name)
8         return tf.Variable(initial)
9
10    @staticmethod
11    def bias_variables(shape, name=None):
12        initial = tf.constant(1.0, shape=shape, name=name)
13        return tf.Variable(initial)
14
15    def __init__(self, session, parent=None):
16        if parent:
17            pass # TODO for GAs some form of deepcopy + modif
18        else:
19            self.initialize_conv1_weights()
20            self.initialize_fc_weights()
21            self.initialize_board_placeholders()
22            self.initialize_training_placeholders()
23            self.initialize_train_vars()
24            self.initialize_convs()
25            self.initialize_ff()
26            self.accum_grads = []
27            self.discount_factor = 0.9
28            self.lambdaa = 0.75
29            self.opt = tf.train.GradientDescentOptimizer(3e-2)
30            self.vars_list = [self.conv1_weights, self.conv1_bias,
31                             self.conv2_weights, self.conv2_bias,
32                             self.conv_diag_weights, self.conv_diag_bias,
33                             self.fc1_weights, self.fc1_bias,
34                             self.fc2_weights,
35                             self.fc2_bias, self.out_weights,
36                             self.out_bias]
37            self.grad_var_list = self.opt.compute_gradients(self.h_out,
38                                                             self.vars_list)
39            self.initialize_accum_grad_vars()
40            self.append_lambda_op()
41            self.values_history = []
```

```

39         session.run(tf.initialize_all_variables())
40
41     # TODO if we ever get to do the evolutionary variant, tic toc...
42     time is pressing
43 @classmethod
44 def spawn(cls, parent):
45     return cls(parent)
46
47 # When we want to start a new match we need to clear some variables
48 def reset_for_game(self):
49     _ = [self.lambda_resets[idx] for idx in
50           range(self.gaccum_list.__len__())]
51     self.values_history = []
52
53 # give an appraisal given a board configuration, train if told to do
54 # so
55 def evaluate(self, boards, session, train=False, verbose=False):
56     boards = np.asarray(boards)
57     if boards.ndim == 2:
58         boards = np.expand_dims(boards, axis=0)
59     batch_size = boards.shape[0]
60     diag = self.get_diagonal(boards)
61     half_sym, chopped = self.turn_boards(boards)
62     if train:
63         v, grad_var_list = session.run([self.h_out,
64                                         self.grad_var_list],
65                                       feed_dict={self.boards_half_sym:
66                                                   half_sym,
67                                                   self.boards_chopped:
68                                                       chopped,
69                                                   self.boards_diag: diag,
70                                                   self.keep_prob: 0.7,
71                                                   self.batch_size:
72                                                       batch_size}) #
73                                                                 self.keep_prob: 1,
74
75         idx = np.argmax(v, 0)
76         self.values_history.append(v[idx])
77         _ = [self.lambda_updates[idx] for idx in
78               range(self.gaccum_list.__len__())]
79     else:
80         v = session.run(self.h_out, feed_dict={self.boards_half_sym:
81                                                 half_sym,
82                                                 self.boards_chopped:
83                                                     chopped,
84                                                 self.boards_diag: diag,
85                                                 self.keep_prob: 1,
86                                                 self.batch_size:
87                                                     batch_size})
88                                                                 #self.keep_prob: 1,
89
90         idx = np.argmax(v, 0)

```

```

76         if verbose:
77             print('max value v at index idx is: ', v[idx][0][0], idx[0])
78         return idx[0], v[idx][0][0] # TODO this looks ugly fix upstream
79
80     # NOTE if training, this should be called at the end of a match
81     def learn_from_outcome(self, tally, session, verbose=False):
82         error = self.values_history[-1] - np.tanh(tally/32)
83         if verbose:
84             print('Loss at end-game: ', loss)
85         for idx, vars in enumerate(self.vars_list):
86             session.run(self.gradient_applications[idx],
87                          feed_dict={self.error: error})
88
89     # TODO fix this, for some reason it's broken on the tf side
90     def set_epochs(self, epochs):
91         self.epochs.assign(epochs)
92
93     def initialize_conv1_weights(self):
94         # create 8 by 1 filter - row/col
95         self.conv1_weights = self.weight_variables([1, 8, 4, 10])
96         self.conv1_bias = self.bias_variables([1, 1, 1, 10])
97
98         # Chop 5x5 part of the board and slide a 3x3 window over it
99         self.conv2_weights = self.weight_variables([3, 3, 8, 10])
100         self.conv2_bias = self.bias_variables([1, 1, 1, 10])
101
102         # a bit of a mis-use of conv-nets but pass the diagonals into a
103         # single input
104         # the useful bits of input channels and features are the reason
105         # for this
106         self.conv_diag_weights = self.weight_variables([1, 1, 4, 4])
107         self.conv_diag_bias = self.bias_variables([1, 1, 1, 4])
108
109     # Try with 64/32 1st layer, second layer neurons
110     def initialize_fc_weights(self):
111         self.fc1_weights = self.weight_variables([174, 64])
112         self.fc1_bias = self.bias_variables([1,64])
113
114         self.fc2_weights = self.weight_variables([64, 32])
115         self.fc2_bias = self.bias_variables([1, 32])
116
117         self.out_weights = self.weight_variables([32,1])
118         self.out_bias = self.bias_variables([1,1])
119
120     def initialize_accum_grad_vars(self):
121         self.gaccum_list = []
122         for gv in self.grad_var_list:
123             g = tf.Variable(tf.zeros(gv[0].get_shape()))
124             self.gaccum_list.append(g)

```

```

123     def append_lambda_op(self):
124         self.lambda_updates = []
125         self.lambda_resets = []
126         self.gradient_applications = []
127         for idx, gv in enumerate(self.grad_var_list):
128             self.lambda_updates.append(
129                 tf.add(self.gaccum_list[idx].__mul__(self.lambdada),
130                     gv[0]))
131             self.lambda_resets.append(self.gaccum_list[idx].__mul__(0))
132             self.gradient_applications.append(
133                 self.opt.apply_gradients(
134                     [(self.gaccum_list[idx].__mul__(self.error), gv[1])]))
135
136     def initialize_train_vars(self):
137         self.epochs = tf.Variable(0)
138
139     def initialize_board_placeholders(self):
140         self.boards_half_sym = tf.placeholder(tf.float32, shape=[None,
141             8, 8, 4])
142         self.boards_chopped = tf.placeholder(tf.float32, shape=[None, 5,
143             8])
144         self.boards_diag = tf.placeholder(tf.float32, shape=[None, 8, 1,
145             4])
146
147     def initialize_training_placeholders(self):
148         self.batch_size = tf.placeholder(tf.int32)
149         self.keep_prob = tf.placeholder(tf.float32)
150         self.error = tf.placeholder(tf.float32)
151
152     @staticmethod
153     def turn_boards(boards):
154         boards = np.asarray(boards)
155         sym1 = boards[:,::-1,:]
156         sym2 = boards[:,:,:-1]
157         sym3 = sym2[:,::-1, :]
158         sym4 = np.transpose(boards, (0, 2, 1))
159         sym5 = np.transpose(sym1, (0, 2, 1))
160         sym6 = np.transpose(sym2, (0, 2, 1))
161         sym7 = np.transpose(sym3, (0, 2, 1))
162         full_sym = np.stack([boards, sym1, sym2, sym3, sym4, sym5, sym6,
163             sym7], axis=3)
164         half_sym = full_sym[:, :, :, 0:4]
165         chopped = full_sym[:, 0:5, 0:5, :]
166         return half_sym, chopped
167
168     @staticmethod
169     def get_diagonal(boards):
170         boards = np.asarray(boards)
171         diags1 = np.expand_dims(boards.diagonal(0, 1, 2), axis=2)

```

```

168     diags2 = np.expand_dims(np.fliplr(boards).diagonal(0,2,1),
169                             axis=2)
170     # Yes, lr does the trick, flipud actually reverses batches
171     # (dimension 0)
172     diags3 = np.fliplr(diags1)
173     diags4 = np.fliplr(diags2)
174     return np.stack([diags1, diags2, diags3, diags4], axis=3)
175
176 def initialize_convs(self):
177     self.h_conv1 = tf.nn.tanh(tf.nn.conv2d(self.boards_half_sym,
178     self.conv1_weights, strides=[1, 1, 1, 1], padding='VALID') +
179     self.conv1_bias)
180
181     self.h_conv2 = tf.nn.tanh(tf.nn.conv2d(self.boards_chopped,
182     self.conv2_weights, strides=[1, 1, 1, 1], padding='VALID') +
183     self.conv2_bias)
184
185     self.h_conv_diag = tf.nn.tanh(tf.nn.conv2d(self.boards_diag,
186     self.conv_diag_weights,
187     strides = [1, 8, 1, 1], padding='VALID') +
188     self.conv_diag_bias)
189
190 def initialize_ff(self):
191     conv1_flat = tf.reshape(self.h_conv1, [-1, 1*8*10])
192     conv2_flat = tf.reshape(self.h_conv2, [-1, 3*3*10])
193     conv_diag_flat = tf.reshape(self.h_conv_diag, [-1, 1*1*4])
194     conv_out = tf.concat(1, [conv1_flat, conv2_flat, conv_diag_flat])
195     self.h_fc1 = tf.nn.tanh(tf.matmul(conv_out,
196     self.fc1_weights)+self.fc1_bias)
197     self.h_fc1_drop = tf.nn.dropout(self.h_fc1, self.keep_prob)
198     self.h_fc2 = tf.nn.tanh(tf.matmul(self.h_fc1_drop,
199     self.fc2_weights) + self.fc2_bias)
200     self.h_fc2_drop = tf.nn.dropout(self.h_fc2, self.keep_prob)
201     self.h_out = tf.nn.tanh(tf.matmul(self.h_fc2_drop,
202     self.out_weights) + self.out_bias)

```


2.b interface

The code was modified in order to accomodate the net playing black or white with no change in its reading of the boards configuration (eg. lines 26, 32 and 64). Printing messages has been suppressed unless calling the interface with a boolean setting 'verbose' (line 31) which is false by default.

```
1  # COMP 6321 Machine Learning, Fall 2016
2  # Federico O'Reilly Regueiro - 40012304
3  # Final project - othello with neural nets
4
5  import othello as o
6  import position as p
7  import board as b
8  import tensorflow as tf
9  import othelloNetV2 as otnet
10 import time
11 import random
12
13 session = tf.Session()
14 game = o.game()
15 on = otnet.othello_net(session)
16 score_series = []
17
18
19 def batch(color):
20     board_now = game.board
21     possible_moves = board_now.get_valid_moves(color)
22     boards = []
23     for m in possible_moves:
24         new_board = b.Board(board_now)
25         new_board.do_move(m, color)
26         new_board.relativize(color)
27         boards.append(new_board.squares)
28     return boards, possible_moves
29
30
31 def play_net(train=False, verbose=False):
32     color = random.choice((b.BLACK, b.WHITE))
33     done = False
34     tic = time.time()
35     global score_series
36     if verbose:
37         print(game.board.to_string())
38     if color == b.WHITE:
39         done = not game.play_random_turn(b.opposite(color), verbose)
40     while not done:
41         boards, moves = batch(color)
42         if boards:
43             if verbose:
```

```

44         print(game.turn_to_string(color), ' plays:')
45         idx, v = on.evaluate(boards, session, train)
46         game.play_move(moves[idx], color)
47         if verbose:
48             print(game.board.to_string())
49         else:
50             game.pass_moves += 1
51         if verbose:
52             print(game.turn_to_string(b.opposite(color)), '\s turn:')
53         done = not game.play_random_turn(b.opposite(color), verbose)
54         outcome = game.board.get_score()
55         color_blind_outcome = {'net': outcome['Black'], 'opponent':
56                               outcome['White']}
57         if train:
58             score_series.append(color_blind_outcome)
59             on.learn_from_outcome(color_blind_outcome['net'] -
60                                 color_blind_outcome['opponent'], session)
61         game.reset()
62         on.reset_for_game()
63         if verbose:
64             print('The round took:', time.time()-tic, ' seconds.')
65             # print('Up to now: ', wins, ' wins, ', losses, ' losses and ',
66                   ties, ' ties.')
67         return color_blind_outcome
68
69     def save_checkpoint(path="./otnet_v2.ckpt"):
70         saver = tf.train.Saver()
71         saver.save(session, path)
72         print('Saved checkpoint')
73
74     def restore_checkpoint(path="./otnet_v2.ckpt"):
75         saver = tf.train.Saver()
76         saver.restore(session, path)
77         print("Model restored.")

```

2.c Driver for training the model, storing the model and a history of scores from simulated matches

The script facilitates training the model while saving both the model and the history of scores every 1000 epochs.

```
1 import pickle
2 import time
3 import othello_interface_v2 as oi2
4 import os.path
5
6 ckpt_fname = 'otnet_v2.ckpt'
7 pckl_fname = 'net_other.pckl'
8
9 if os.path.isfile('otnet_v2.ckpt'):
10     oi2.restore_checkpoint()
11     if os.path.isfile(pckl_fname):
12         f = open(pckl_fname, 'rb')
13         x = pickle.load(f)
14         f.close()
15         oi2.score_series = x
16
17 for batch in range(200):
18     tic = time.time()
19     for _ in range(1000):
20         oi2.play_net(True)
21     print('batch ', batch, ' took ', (time.time()-tic)/60, ' minutes')
22     tl = time.localtime()
23     print('finished at: ', tl.tm_hour, ':', tl.tm_min)
24     oi2.save_checkpoint()
25     oi2.save_checkpoint('SavedModels/otnetv2'+str(oi2.score_series.__len__())+'.ckpt')
26     f = open(pckl_fname, 'wb')
27     pickle.dump(oi2.score_series, f)
28     f.close()
```

2.d Displaying results

Simple routine for loading and displaying results in another python interpreter while the model is still training.

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import pickle
4 import os
5
6
7 pckl_fname = 'net_other.pckl'
8
9
10 def moving_average(a, n=3):
11     ret = np.cumsum(a, dtype=float)
12     ret[n:] = ret[n:] - ret[:-n]
13     return ret[n-1:]/n
14
15
16 def load_and_display():
17     if not os.path.isfile(pckl_fname):
18         raise FileNotFoundError
19     f = open(pckl_fname, 'rb')
20     x = pickle.load(f)
21     f.close()
22     n = []
23     o = []
24     for sc in x:
25         n.append(sc['net'])
26         o.append(sc['opponent'])
27
28     n = np.asarray(n)
29     o = np.asarray(o)
30
31     window_length = n.__len__() / 10
32     mav = moving_average((n-o), window_length)
33     plt.plot(mav, 'b', np.zeros(mav.__len__()), 'k--')
34     plt.title('scores-moving average, rectangular window,
35              n='+str(window_length))
36     plt.xlabel('epochs')
37     plt.ylabel('score moving average')
38     plt.show()
39
40     mav2 = moving_average((n-o) > 0, window_length)
41     plt.plot(mav2, 'b', 0.5 * np.ones(mav2.__len__()), 'k--')
42     plt.title('games won moving average, rectangular window,
43              n='+str(window_length))
44     plt.xlabel('epochs')
45     plt.ylabel('wins moving average')
```

```

44 plt.show()
45
46 # plt.plot(np.cumsum((n-o) > 0), 'g', np.cumsum((n-o) == 0), 'b',
      np.cumsum((n-o) < 0), 'r')
47 # plt.legend(['wins', 'ties', 'losses'], loc='upper left')
48 # plt.title('Progression of outcomes over epochs')
49 # plt.xlabel('epochs')
50 # plt.ylabel('accumulated outcomes')
51 # plt.show()
52
53 plt.plot(np.divide((np.cumsum((n - o) > 0))+1, np.cumsum((n-o) <
      0)+1), 'g', np.ones(n.__len__()), 'k--')
54 plt.ylim([0.85, 1.02])
55 plt.title('Ratio of wins over losses over epochs')
56 plt.xlabel('epochs')
57 plt.ylabel('wins to losses')
58 plt.show()

```

References

- [1] Kumar Chellapilla and David B Fogel. Evolution, neural networks, games, and intelligence. *Proceedings of the IEEE*, 87(9):1471–1496, 1999.
- [2] Siang Y. Chong, Mei K. Tan, and Jonathon D. White. Observing the evolution of neural networks learning to play the game of othello. *IEEE Transactions on Evolutionary Computation*, 9:240–251, 2005.
- [3] Anton V. Leouski and Paul E. Utgoff. What a neural network can learn about othello, 1996.
- [4] Gerald Tesauro. Practical issues in temporal difference learning. In *Machine Learning*, pages 257–277, 1992.