

Comp 6321 - Machine Learning

Using Neural Nets for playing othello

Federico O'Reilly Regueiro

Concordia University

November 30, 2016

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - One of the oldest questions in AI
 - The trick is in finding ways to narrow the search
 - Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - One of the oldest questions in AI
 - The trick is in finding ways to narrow the search
 - Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - One of the oldest questions in AI
 - The trick is in finding ways to narrow the search
 - Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - One of the oldest questions in AI
 - The trick is in finding ways to narrow the search
 - Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - ▶ One of the oldest questions in AI
 - ▶ The trick is in finding ways to narrow the search
 - ▶ Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - ▶ One of the oldest questions in AI
 - ▶ The trick is in finding ways to narrow the search
 - ▶ Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - ▶ One of the oldest questions in AI
 - ▶ The trick is in finding ways to narrow the search
 - ▶ Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - ▶ One of the oldest questions in AI
 - ▶ The trick is in finding ways to narrow the search
 - ▶ Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - ▶ One of the oldest questions in AI
 - ▶ The trick is in finding ways to narrow the search
 - ▶ Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ▶ ...the feedback becomes very sparse

Problem statement

- Zero-sum, perfect-knowledge (no chance involved) competitive-game
 - ▶ A sandbox - toy-representation of reality
 - ▶ bounded problem space with clear goal and set of rules
 - ▶ bounded, but can be huge (ie, GO - 10^{761} possible games!) [1]
- Can a machine learn to play
 - ▶ One of the oldest questions in AI
 - ▶ The trick is in finding ways to narrow the search
 - ▶ Has been well answered, requiring less expert knowledge each time
- Without expert knowledge (rules or labels)...
 - ▶ ...the feedback becomes very sparse

Approaches and solutions - common elements

- Classification problem

- ▶ Dual class

- given a game-state, what are the odds of winning

- Learn a policy for action given a state $P(a|s)$

- ▶ Multi-class

- given a game-state, what is the best next move

- Learn a policy for action given a state $P(a|s)$

Approaches and solutions - common elements

- Classification problem

- ▶ Dual class

- given a game-state, what are the odds of winning

- ★ Look ahead n-moves - (n-ply) - then decide best path given leaf 'value'

- ▶ Multi-class

- given a game-state, what is the best next move

- Learn a policy for action given a state $\pi(a|s)$

Approaches and solutions - common elements

- Classification problem

- ▶ Dual class

- given a game-state, what are the odds of winning

- Look ahead n -moves - (n -ply) - then decide best path given leaf 'value'

- ▶ Multi-class

- given a game-state, what is the best next move

- ★ Learn a 'policy' for action given a state $P(a|s)$

Approaches and solutions - common elements

- Classification problem

- ▶ Dual class

- given a game-state, what are the odds of winning

- ★ Look ahead n -moves - (n -ply) - then decide best path given leaf 'value'

- ▶ Multi-class

- given a game-state, what is the best next move

- Learn a 'policy' for action given a state $P(a|s)$

Approaches and solutions - common elements

- Classification problem

- ▶ Dual class

- given a game-state, what are the odds of winning

- Look ahead n -moves - (n -ply) - then decide best path given leaf 'value'

- ▶ Multi-class

- given a game-state, what is the best next move

- ★ Learn a 'policy' for action given a state $P(a|s)$

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge

- ▶ e.g. Deep Blue

- Supervised learning - collect labeled states and train

- Also depends on human expert knowledge

- Labor intensive collection and labeling

- Genetic optimizations - Evolutionary NNs

- Does not exploit NNs learning capabilities
but won't get stuck on local minima...

- Slow to converge

- Capable of finding innovative strategies [4] [2]

- Reinforcement learning

- TD-learning

- Lecture 7: TD-Reinforcement Learning

- Like having sparse and time-delayed labels

- Credit assignment problem

- explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue

- Supervised learning - collect labeled states and train

Also depends on human expert knowledge

Labor intensive collection and labeling

- Genetic optimizations - Evolutionary NNs

Does not exploit NNs learning capabilities
but won't get stuck on local minima...

Slow to converge

Capable of finding innovative strategies [4] [2]

- Reinforcement learning

TD-learning

John Sutton's TD-Backgammon - classic example

Like having sparse and time-delayed labels

Credit assignment problem

explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - Slow to converge
 - Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - TD-learning
 - Like having sparse and time-delayed labels
 - Credit assignment problem
 - explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - Slow to converge
 - Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - TD-learning
 - Like having sparse and time-delayed labels
 - Credit assignment problem
 - explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - Slow to converge
 - Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - TD-learning
 - Value Iteration, TD-Error, Monte Carlo
 - Like having sparse and time-delayed labels
 - Credit assignment problem
 - explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ▶ Tesauro's TD-Backgammon - chance element
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ★ Tesauro's TD-Backgammon - chance element
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ★ Tesauro's TD-Backgammon - chance element
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ★ Tesauro's TD-Backgammon - chance element
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ★ Tesauro's TD-Backgammon - chance element
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Approaches and solutions to the problem

- Rule-based approach dependent on expert knowledge
 - ▶ e.g. Deep Blue
- Supervised learning - collect labeled states and train
 - ▶ Also depends on human expert knowledge
 - ▶ Labor intensive collection and labeling
- Genetic optimizations - Evolutionary NNs
 - ▶ Does not exploit NNs learning capabilities but won't get stuck on local minima...
 - ▶ Slow to converge
 - ▶ Capable of finding innovative strategies [4] [2]
- Reinforcement learning
 - ▶ TD-learning
 - ★ Tesauro's TD-Backgammon - chance element
 - ▶ Like having sparse and time-delayed labels
 - ▶ Credit assignment problem
 - ▶ explore-exploit dilemma

Current state-of-the-art

- Alpha-go
 - ▶ Two policy convolutional networks - 1 large, 1 small - prune search tree $TD(\lambda)$
 - ▶ One Fully connected - predict win value
- DeepMind Atari deep reinforcement learning
 - ▶ Deep neural nets meet $TD(\lambda)$

Current state-of-the-art

- Alpha-go
 - ▶ Two policy convolutional networks - 1 large, 1 small - prune search tree $TD(\lambda)$
 - ▶ One Fully connected - predict win value
- DeepMind Atari deep reinforcement learning
 - Deep neural nets meet $TD(\lambda)$

Current state-of-the-art

- Alpha-go
 - ▶ Two policy convolutional networks - 1 large, 1 small - prune search tree $TD(\lambda)$
 - ▶ One Fully connected - predict win value
- DeepMind Atari deep reinforcement learning
 - Deep neural nets meet $TD(\lambda)$

Current state-of-the-art

- Alpha-go
 - ▶ Two policy convolutional networks - 1 large, 1 small - prune search tree $TD(\lambda)$
 - ▶ One Fully connected - predict win value
- DeepMind Atari deep reinforcement learning
 - ▶ Deep neural nets meet $TD(\lambda)$

Current state-of-the-art

- Alpha-go
 - ▶ Two policy convolutional networks - 1 large, 1 small - prune search tree $TD(\lambda)$
 - ▶ One Fully connected - predict win value
- DeepMind Atari deep reinforcement learning
 - ▶ Deep neural nets meet $TD(\lambda)$

Current project state - discarded approaches

- Supervised Learning

- ▶ Acquiring sets is a cumbersome task
requires an overhead outside of ML - eg Edax

- Rule-based

Heuristic - Decision tree - in place but focus is on nets

Current project state - discarded approaches

- Supervised Learning
 - ▶ Acquiring sets is a cumbersome task
requires an overhead outside of ML - eg Edax
- Rule-based
 - Heuristic - Decision tree - in place but focus is on nets

Current project state - discarded approaches

- Supervised Learning

- ▶ Acquiring sets is a cumbersome task
requires an overhead outside of ML - eg Edax

- Rule-based

- ▶ Heuristic - Decision tree - in place but focus is on nets

Current project state - discarded approaches

- Supervised Learning

- ▶ Acquiring sets is a cumbersome task
requires an overhead outside of ML - eg Edax

- Rule-based

- ▶ Heuristic - Decision tree - in place but focus is on nets

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

Based on Chelapilla and Fogel[2]

Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.

$$\sigma'_i(j) = \sigma_i(j) \exp(r N_j(0, 1))$$

$$w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

Based on Chelapilla and Fogel[2]

Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.

$$\sigma'_i(j) = \sigma_i(j) \exp(r N_j(0, 1))$$

$$w'_i(j) = w_i(j) + \sigma'_i(j) N_j(0, 1)$$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

Based on Chelapilla and Fogel[2]

Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.

$$\sigma'_i(j) = \sigma_i(j) \exp(r N_j(0, 1))$$

$$w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

Based on Chelapilla and Fogel[2]

Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.

$$\sigma'_i(j) = \sigma_i(j) \exp(r N_j(0, 1))$$

$$w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

Based on Chelapilla and Fogel[2]

Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.

$$\sigma'_i(j) = \sigma_i(j) \exp(r N_j(0, 1))$$

$$w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

- ▶ Based on Chelapilla and Fogel[2]
- ▶ Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.
- ▶ $\sigma'_i(j) = \sigma_i(j) \exp(\tau N_j(0, 1))$
- ▶ $w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

- ▶ Based on Chelapilla and Fogel[2]
- ▶ Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.
- ▶ $\sigma'_i(j) = \sigma_i(j) \exp(\tau N_j(0, 1))$
- ▶ $w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

- ▶ Based on Chelapilla and Fogel[2]
- ▶ Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.
- ▶ $\sigma'_i(j) = \sigma_i(j) \exp(\tau N_j(0, 1))$
- ▶ $w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

- ▶ Based on Chelapilla and Fogel[2]
- ▶ Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.
- ▶ $\sigma'_i(j) = \sigma_i(j) \exp(\tau N_j(0, 1))$
- ▶ $w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

- ▶ Based on Chelapilla and Fogel[2]
- ▶ Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.
- ▶ $\sigma'_i(j) = \sigma_i(j) \exp(\tau N_j(0, 1))$
- ▶ $w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$

Current project state - narrowed approaches

- TD learning

- ▶ Similar to back propagation but recurses temporally
- ▶ $\Delta_{w_t} = \alpha(P_{t+1} - P_t) \sum_{k=1}^t \lambda^{t-k} \nabla_w P_k, \quad 0 \leq \lambda \leq 1$
- ▶ Based on Leouski and Utgoff's paper[3]
- ▶ They use symmetry, rotation and weight sharing - 96 h.u.
- turn into conv net

- ENN

- ▶ Based on Chelapilla and Fogel[2]
- ▶ Generation has 15 strategies, change vector $\sigma_i(j)$ for j^{th} weight of i^{th} strategy.
- ▶ $\sigma'_i(j) = \sigma_i(j) \exp(\tau N_j(0, 1))$
- ▶ $w'_i(j) = w_i(j) + \sigma_i(j) N_j(0, 1)$

References



Christopher Burger.

Google deepmind's alphago: How it works, 2016.



Kumar Chellapilla and David B Fogel.

Evolution, neural networks, games, and intelligence.

Proceedings of the IEEE, 87(9):1471–1496, 1999.



Anton V. Leouski and Paul E. Utgoff.

What a neural network can learn about othello, 1996.



David Moriarty and Risto Miikkulainen.

Evolving complex othello strategies using marker-based genetic encoding of neural networks.

Technical report, 1993.